

Choice of Variables in Vector Autoregressions*

Marek Jarociński

Bartosz Maćkowiak

European Central Bank

European Central Bank and CEPR

September 2011

Abstract

Suppose that a dataset with N time series is available. $N_1 < N$ of those are the variables of interest. You want to estimate a vector autoregression (VAR) with the variables of interest. Which of the remaining $N - N_1$ variables, if any, should you include in the VAR with the variables of interest? We develop a Bayesian methodology to answer this question. This question arises in most applications of VARs, whether in forecasting or impulse response analysis. We apply the methodology to the euro area data and find that when the variables of interest are the price level, GDP, and the short-term interest rate, the VAR with these variables should also include the unemployment rate, the spread between corporate bonds and government bonds, the purchasing managers index, and the federal funds rate. Of independent interest, we develop Bayesian tests of block-exogeneity – Granger causality – in VARs.

Keywords: Bayesian vector autoregression, Bayesian model choice, block-exogeneity, Granger causality. (*JEL:* C32, C52, C53.)

*Jarociński: European Central Bank, Kaiserstrasse 29, 60311 Frankfurt am Main, Germany (e-mail: marek.jarocinski@ecb.int); Maćkowiak: European Central Bank, Kaiserstrasse 29, 60311 Frankfurt am Main, Germany (e-mail: bartosz.mackowiak@ecb.int). We thank for helpful comments Gianni Amisano, John Geweke, David Madigan, Tao Zha and especially Chris Sims. We also thank seminar participants at the National Bank of Poland and Humboldt University Berlin. The views expressed in this paper are solely those of the authors and do not necessarily reflect the views of the European Central Bank.

1 Introduction

Vector autoregressions (VARs) are a standard tool for forecasting and impulse response analysis. When we set out to estimate a VAR, we rarely know a priori which variables to include in the VAR. Typically, we have a small number of variables we are interested in and we know that, in principle, we can include many other variables in the VAR. For example, when we want to forecast inflation and GDP, we realize that many variables may improve our forecast. To take another example, when we want to estimate the impulse response of hours worked to a technology shock, we realize that the inclusion or exclusion of many “control variables” may affect our estimate. Often, we include or exclude one variable at a time and evaluate informally whether “the results change”.

This paper studies formally the choice of variables in VARs. We study the following question. Suppose that a dataset with N time series is available. $N_1 < N$ of these time series are *the variables of interest*. Let all other $N - N_1$ time series be called *the remaining variables*. We want to estimate a VAR with the variables of interest. Which of the remaining variables, if any, should we include in the VAR with the variables of interest?

We develop a methodology to answer this question. The idea behind this methodology is as follows. Let y be the vector of all variables in the dataset. Consider a partition of y into two subvectors, $y = \{y_i, y_j\}$, such that all variables of interest are elements of y_i . Note that y_i includes all variables of interest and possibly some remaining variables and y_j includes the other remaining variables. Suppose that y_i is block-exogenous to y_j . The statement “ y_i is block-exogenous to y_j ” means that “the elements of y_j are of no help in improving a forecast of any variable contained in y_i based on lagged values of all the elements of y_i alone”.¹ If y_i is block-exogenous to y_j , one can exclude y_j from a VAR model with y_i . In other words, only those remaining variables that are elements of y_i are helpful in modeling the variables of interest. This reasoning leads to the following conclusion. The decision about which variables to include in a VAR with variables of interest involves evaluating block-exogeneity restrictions in a VAR with all variables. The block-exogeneity restrictions have the form:

¹The statement “ y_i is block-exogenous to y_j ” is the same as the statement that “the variables in y_j do not Granger-cause any of the variables in y_i ”. See Section 2.1. Throughout this paper we use the term “block-exogeneity”, but we could have also used the term “Granger causality”.

“the variables of interest and a subset of the remaining variables are block-exogenous to every other variable”.

We develop the Bayesian implementation of this methodology. We work with Bayesian VARs, because Bayesian VARs are popular and we find Bayesian inference appealing. To evaluate a block-exogeneity restriction in a VAR, a Bayesian compares the marginal likelihood of the data implied by that VAR without restrictions with the marginal likelihood of the data implied by that VAR given the block-exogeneity restriction. It is well known that, with conjugate priors, the marginal likelihood implied by an unrestricted VAR can be evaluated analytically. We show how to evaluate conveniently the marginal likelihood implied by a VAR with one block-exogeneity restriction and multiple block-exogeneity restrictions. The marginal likelihood implied by a VAR with one block-exogeneity restriction can be evaluated analytically. The marginal likelihood implied by a VAR with multiple block-exogeneity restrictions can be evaluated with a simple Monte Carlo. Given these results, a Bayesian can reach a decision about which variables to include in a VAR with variables of interest and he or she can do so quickly.

The results that we work out rely on the assumption that the prior in each restricted VAR has to be consistent with the prior in the unrestricted VAR, in a sense that we will make precise.² This assumption is simple and natural. Furthermore, the prior in the unrestricted VAR is assumed to be conjugate. In an application, we use the standard prior employed in Bayesian VARs due to Sims and Zha (1998).

The output of our methodology are posterior probabilities of models. Therefore, after implementing our methodology one can in principle compute forecasts or impulse responses by averaging across models. We prefer to use a single VAR after implementing our methodology, the VAR with maximum posterior probability, because we want to treat this VAR as a benchmark to be used in future research, including for comparison with models more complex than VARs. Formally, our focus on the VAR with maximum posterior probability is justified under a zero-one loss function.

Alternatives. We discuss in the paper several alternatives to the methodology that we propose. The alternatives that we discuss compare models based on a predictive density

²See Section 2.2.

of the variables of interest. In contrast, our methodology compares models based on the marginal likelihood – prior predictive density – of all variables in the dataset. We argue that the alternatives are less attractive than our methodology. For example, we explain why it is unappealing to use as a criterion for the choice of variables the prior predictive density of the variables of interest, that is the marginal likelihood marginalized with respect to any remaining variables.³

The following approach to answer the question that we study is popular: (i) consider a family of VARs such that each VAR includes the variables of interest and a subset of the remaining variables, (ii) for each VAR, compute root mean squared errors of out-of-sample point forecasts of the variables of interest, and (iii) declare the VAR with smallest root mean squared errors as the best VAR with the variables of interest. Computation of root mean squared errors of point forecasts as a means of model choice is a widespread practice also in work with Bayesian VARs. We find this practice unappealing, because this practice relies on point forecasts and disregards the uncertainty of the forecasts. We propose a convenient alternative with a formal Bayesian justification.

The results of independent interest. The results concerning how to evaluate conveniently the marginal likelihood implied by a VAR with block-exogeneity are of independent interest. Since Granger (1969) and Sims (1972), there has been a significant interest in testing Granger causality and, since Sims (1980), often in VARs. Tests of block-exogeneity have been performed using the frequentist likelihood ratio test, even in Bayesian VARs, or the Schwarz criterion which is only an asymptotic approximation to a Bayesian test.⁴ The properties of the likelihood ratio test of a zero restriction in a Bayesian VAR with the standard Sims-Zha prior are unclear, given that this prior shrinks coefficients to zero. Formal Bayesian tests have been possible in principle, though essentially unused in practice because they require cumbersome Monte Carlo.⁵ We show that the marginal likelihood implied by a

³See Section 4.

⁴For example, Cushman and Zha (1997) use the likelihood ratio test and Maćkowiak (2007) uses the Schwarz criterion. Both papers use Bayesian VARs.

⁵For example, one can use the Gibbs sampler developed in Waggoner and Zha (2003) to sample from the posterior density of the parameters of a VAR with block-exogeneity and then use the method of Chib (1995) to compute from the Gibbs output the marginal likelihood implied by that VAR.

VAR with one block-exogeneity restriction can be evaluated analytically; furthermore, the marginal likelihood implied by a VAR with multiple block-exogeneity restrictions can be evaluated with a simple Monte Carlo. We recommend that Bayesians use these results also when their interest is different than the choice of variables in a VAR.

Application. We apply our methodology to study the following question. We want to estimate a VAR with a measure of the price level, GDP, and a short-term interest rate controlled by monetary policy in the euro area. Which other macroeconomic and financial variables, if any, should we include in that VAR? We consider a quarterly dataset with twenty variables, of which three are the variables of interest (the Harmonized Index of Consumer Prices, real GDP, and the overnight interbank interest rate Eonia) and seventeen macroeconomic and financial variables are the remaining variables. We find that the best VAR with the price level, GDP, and the policy rate includes in addition the following four variables: the unemployment rate, the spread between corporate bonds and government bonds of identical maturity, the purchasing managers index, and the federal funds rate. Thus if one wants to model the price level, GDP, and the policy rate in the euro area in a VAR, that VAR should also include a measure of capacity utilization, a notion central to Keynesian and New Keynesian business cycle models (the unemployment rate), a variable central to business cycle models with financial frictions (the bond spread), the main leading indicator (the purchasing managers index), and a variable external to the euro area (the federal funds rate). We think that this is a plausible finding.

Contacts with existing work. Bańbura et al. (2010) show that a large Bayesian VAR with as many as 131 variables yields smaller root mean squared errors of out-of-sample point forecasts of a few variables of interest compared with small VARs. This finding appears to suggest that one can simply estimate a VAR with all variables in one's dataset. However, Bańbura et al. (2010) also show that a VAR with 20 variables achieves much of the improvement in the predictive performance over small VARs. This finding raises the following questions: How do we decide which 20 variables to include in a VAR? How do we decide whether the optimal number of variables to include is 20 or some other number? Our methodology addresses these questions in a systematic way. Whether a small, medium-size, or large model is preferred is in general sample-specific.

Zha (1999) and Waggoner and Zha (2003) develop Monte Carlo methods for Bayesian inference concerning parameters of a VAR with block-exogeneity. In contrast, we are interested in evaluating the marginal likelihood implied by a VAR with block-exogeneity and we are not interested in inference concerning parameters of a VAR with block-exogeneity. In the applications of VARs with block-exogeneity, Cushman and Zha (1997) and Zha (1999) are interested in all variables being modeled and do not use block-exogeneity to justify dropping variables. Furthermore, in this work block-exogeneity is either imposed or tested from the frequentist perspective.

Similar to us, George et al. (2008) study Bayesian VARs with zero restrictions. However, the zero restrictions in George et al. (2008) are a priori independent across individual coefficients. In contrast, we are concerned with zero restrictions that are not independent across individual coefficients and instead always apply to blocks of coefficients. Second, George et al. (2008) aim to do inference in the VAR with all variables, averaging over different possible restrictions. In contrast, we pick the single best restriction, because our goal is to choose the optimal benchmark VAR of a reduced dimension. The numerical methods of George et al. (2008) are not suitable for picking the single best restriction or evaluating posterior probabilities of restrictions.

Outline of the rest of this paper. Section 2 states the question that we study in this paper and proposes a methodology to answer this question. Section 3 describes how to implement the methodology that we propose. Section 4 discusses some alternative methodologies to answer the question that we study. Section 5 contains the application to the euro area data. Section 6 concludes. There are two appendices, one with details concerning the prior used in Section 5 and one with details concerning the algorithm used in Section 5. In addition, a Technical Appendix is available.

2 Question and methodology to answer it

This section states the question that we study in this paper and proposes a methodology to answer this question.

Throughout this paper, we consider VAR models all of which have the form

$$y(t) = \gamma + B(L)y(t-1) + u(t), \quad (1)$$

where $y(t)$ is a vector of variables observed in period $t = 1, \dots, T$, γ is a constant term, $B(L)$ is a matrix polynomial in the lag operator of order $P - 1$, and $u(t)$ is a Gaussian vector with mean zero and variance-covariance matrix Σ , conditional on $y(t-s)$ for all $s \geq 1$.

The question that we study in this paper is the following. Suppose that a dataset with $T + P$ observations of N variables is available. Furthermore, suppose that $N_1 < N$ of those variables are *the variables of interest*. We refer to all other N_2 variables in the dataset, where $N_2 = N - N_1$, as *the remaining variables*. We want to estimate a VAR with the variables of interest. Which of the remaining variables, if any, should we include in the VAR with the variables of interest?

The rest of this section proposes a methodology to answer this question. The idea is to cast the choice of variables in a VAR as the choice of one model from the family of VARs with block-exogeneity.

2.1 Block-exogeneity

We begin by defining block-exogeneity and making several observations about it.

Consider a partition of a vector y into two subvectors, $y = \{y_i, y_j\}$, and a conformable partition of the VAR model of y

$$\begin{pmatrix} y_i(t) \\ y_j(t) \end{pmatrix} = \gamma + \begin{pmatrix} B_{ii}(L) & B_{ij}(L) \\ B_{ji}(L) & B_{jj}(L) \end{pmatrix} \begin{pmatrix} y_i(t-1) \\ y_j(t-1) \end{pmatrix} + u(t). \quad (2)$$

Definition 1 *Block-exogeneity*: *The vector of variables y_i is said to be block-exogenous with respect to the vector of variables y_j if the elements of y_j are of no help in improving a forecast of any variable contained in y_i that is based on lagged values of all the elements of y_i alone.*⁶

Relation to a zero restriction: In the VAR given in equation (2), y_i is block-exogenous to y_j if and only if $B_{ij}(L) = 0$.

⁶This definition is taken from Hamilton (1994), p.309.

Relation to Granger causality: The restriction $B_{ij}(L) = 0$ is the same as the statement that the variables in y_j do not Granger-cause any of the variables in y_i .⁷ Throughout this paper we use the term “block-exogeneity”, but we could have also used the term “Granger causality”.

Relation to block-recursive form: With the restriction $B_{ij}(L) = 0$ the VAR given in equation (2) is block-recursive and has two blocks. In the first block, current y_i is explained only by lagged values of itself; and in the second block, current y_j is explained by lagged values of itself and lagged values of y_i . This means that y_j is unhelpful if the interest is in modeling some or all elements of y_i . Only the first block is relevant if the interest is in modeling some or all elements of y_i . Suppose that all variables of interest are elements of y_i . Then y_j is unhelpful in modeling the variables of interest. Only the elements of y_i are relevant for modeling the variables of interest.

Bayesian test of block-exogeneity: Let $p(Y)$ denote the marginal likelihood of the data implied by the VAR given in equation (2). Let $p(Y|B_{ij}(L) = 0)$ denote the marginal likelihood of the data implied by that VAR conditional on the restriction $B_{ij}(L) = 0$. This paper uses Bayesian inference. To evaluate the restriction $B_{ij}(L) = 0$, a Bayesian compares the marginal likelihood implied by the unrestricted VAR with the marginal likelihood implied by the restricted VAR. When the prior probability of the unrestricted model is equal to the prior probability of the restricted model, the Bayesian prefers the model implying a higher marginal likelihood. The Bayes factor is defined as

$$\frac{p(Y|B_{ij}(L) = 0)}{p(Y)}.$$

With equal prior probabilities, the Bayes factor equals the posterior odds in favor of the restriction.

These observations suggest the following principle that we adopt. The Bayesian decision about which variables to include in a VAR with variables of interest involves evaluating, via marginal likelihood, block-exogeneity restrictions in a VAR with all variables. The block-exogeneity restrictions have the form: “the variables of interest and a subset of the

⁷See Hamilton (1994), p.303, for the definition of non-Granger causality. The restriction $B_{ij}(L) = 0$ is also the same as the statement that the variables in y_i are Granger causally prior to the variables in y_j . See Sims (2010) for the definition of Granger causal priority.

remaining variables are block-exogenous to every other variable”.

2.2 A family of VARs and the choice of the best VAR

We now state our methodology formally. Namely, we define the family of VARs with block-exogeneity and we cast the choice of variables in a VAR as the choice of one model from this family.

Let y denote the vector consisting of all N variables in the dataset. Consider a partition of y into two subvectors, $y = \{y_1, y_2\}$, where y_1 denotes N_1 variables of interest and y_2 denotes N_2 remaining variables. Let Ω denote a family of models such that: (i) each model in the family Ω is a VAR with N variables represented by y with zero, one, or more block-exogeneity restrictions, and (ii) in each model in the family Ω , the variables represented by y_1 are in the first block of that model.⁸

Let $p(B, \Sigma | \omega^U)$ denote the prior density of B and Σ in the unrestricted model $\omega^U \in \Omega$, where B is a matrix that collects from equation (1) the parameters in the matrix polynomial $B(L)$ and the constant term γ . Furthermore, for any restricted model $\omega^R \in \Omega$, let $p(B_U, \Sigma | \omega^R)$ denote the prior density of B_U – the unrestricted elements of B – and Σ in that model. We assume that the prior density $p(B_U, \Sigma | \omega^R)$ satisfies the following property

$$p(B_U, \Sigma | \omega^R) = p(B_U, \Sigma | \omega^U, B_R = 0), \quad (3)$$

where B_R denotes the elements of B that are set to zero reflecting a single block-exogeneity restriction or multiple block-exogeneity restrictions in the model ω^R . We find assumption (3) simple and natural. If we think that the prior density in the unrestricted model $\omega^U \in \Omega$ is $p(B, \Sigma | \omega^U)$, then it is reasonable to think that the prior density in a given restricted model $\omega^R \in \Omega$ is equal to the prior density in the unrestricted model ω^U conditional on the restriction in the model ω^R . This is what assumption (3) states. Essentially, the prior in each restricted VAR has to be consistent with the prior in the unrestricted VAR. Note that assumption (3) is important for the results in Section 3.3.

⁸A VAR with zero block-exogeneity restrictions has one block. Furthermore, the family Ω contains one unrestricted model and a number of restricted models, where the number of restricted models depends on N_2 . In Section 3.4, we state how many restricted models there are in the family Ω for each value of N_2 .

We assume that all models in the family Ω have equal prior probabilities. Under this assumption, posterior odds between any two models in the family Ω are equal to the ratio of the marginal likelihoods of the data implied by these two models. In the rest of this paper, we focus on marginal likelihoods. It is trivial to adapt our methodology to the case when models in the family Ω have different prior probabilities.

The methodology that we propose can be stated as follows: Evaluate the marginal likelihood of the data implied by each model in the family Ω , that is evaluate $p(Y|\omega)$ for each $\omega \in \Omega$, and choose the model associated with the highest marginal likelihood, ω^* .

In the end, the model of interest is only the first block of the model ω^* . That is, having implemented this methodology, we use only the first block of the model ω^* for forecasting or impulse response analysis. The reasons are that: (i) the vector of the variables of interest, y_1 , is in the first block of the model ω^* , and (ii) once we have found the best model ω^* , given the definition of block-exogeneity, only the first block of that model is relevant for modeling y_1 .

2.3 Example

Suppose that $N = 3$ and $N_1 = 1$, i.e. the dataset contains three variables and there is one variable of interest. Setting $N = 3$ we can rewrite equation (1) as follows

$$\begin{pmatrix} y_1(t) \\ y_{2,1}(t) \\ y_{2,2}(t) \end{pmatrix} = \gamma + \begin{pmatrix} B_{11}(L) & B_{12}(L) & B_{13}(L) \\ B_{21}(L) & B_{22}(L) & B_{23}(L) \\ B_{31}(L) & B_{32}(L) & B_{33}(L) \end{pmatrix} \begin{pmatrix} y_1(t-1) \\ y_{2,1}(t-1) \\ y_{2,2}(t-1) \end{pmatrix} + u(t). \quad (4)$$

We think of y_1 as the variable of interest and we think of $y_{2,1}$ and $y_{2,2}$ as the remaining variables. We study the following question: In addition to y_1 , should we include in the VAR $y_{2,1}$ and $y_{2,2}$, only $y_{2,1}$, only $y_{2,2}$, or neither $y_{2,1}$ nor $y_{2,2}$?

We make four observations about this example.

First, “including $y_{2,1}$ and $y_{2,2}$ in addition to y_1 ” means estimating the unrestricted VAR given in equation (4). “Including only $y_{2,1}$ in addition to y_1 ” means estimating the VAR given in equation (4) with the restriction $B_{13}(L) = B_{23}(L) = 0$, i.e. with the restriction that y_1 and $y_{2,1}$ are block-exogenous to $y_{2,2}$. “Including only $y_{2,2}$ in addition to y_1 ” means estimating that VAR with the restriction $B_{12}(L) = B_{32}(L) = 0$, i.e. with the restriction

that that y_1 and $y_{2,2}$ are block-exogenous to $y_{2,1}$. “Including neither $y_{2,1}$ nor $y_{2,2}$ in addition to y_1 ” means estimating that VAR with the restriction $B_{12}(L) = B_{13}(L) = 0$, i.e. with the restriction that y_1 is block-exogenous to $y_{2,1}$ and $y_{2,2}$.

Second, with each block-exogeneity restriction the VAR given in equation (4) is block-recursive. For example, consider the restriction $B_{13}(L) = B_{23}(L) = 0$. With this restriction, current y_1 and $y_{2,1}$ are explained only by lagged values of themselves, and current $y_{2,2}$ is explained by lagged values of itself and lagged values of y_1 and $y_{2,1}$. This means that $y_{2,1}$ is helpful and $y_{2,2}$ is unhelpful when the interest is in modeling y_1 .⁹

Third, to test the block-exogeneity restrictions a Bayesian evaluates the marginal likelihood of the data implied by the VAR given in equation (4) without and with each block-exogeneity restriction. When the prior probabilities of the unrestricted model and the restricted models are equal, the Bayesian prefers the specification associated with the highest marginal likelihood.

Fourth, the following discussion explains why we also consider models with more than one block-exogeneity restriction.¹⁰ We do so because of the concern that the result of the block-exogeneity tests discussed so far depends on whether some other restriction is imposed. In particular, one could be concerned that modeling the interaction between the remaining variables matters for the result of the block-exogeneity tests discussed so far. Our approach to modeling the interaction between the remaining variables is to consider block-exogeneity restrictions among the remaining variables. This approach leads us to consider two additional models: (i) the VAR given in equation (4) with the restriction $B_{12}(L) = B_{13}(L) = B_{23}(L) = 0$, i.e. with the restriction that y_1 is block-exogenous to $y_{2,1}$ and $y_{2,2}$ and $y_{2,1}$ is block-exogenous to $y_{2,2}$, and (ii) the VAR given in equation (4) with the restriction $B_{12}(L) = B_{13}(L) = B_{32}(L) = 0$, i.e. with the restriction that y_1 is block-exogenous to $y_{2,1}$ and $y_{2,2}$ and $y_{2,2}$ is block-exogenous to $y_{2,1}$. Note that with either of those two restrictions, the decision is “include neither $y_{2,1}$ nor $y_{2,2}$ in addition to y_1 ”.

⁹We use the terms “block-exogeneity” and “block-recursive” for consistency throughout the paper, even though in the example studied in this section some “blocks” include only one variable.

¹⁰In other words, the following discussion explains why we also consider block-recursive models with more than two blocks.

3 Methodology: implementation

This section describes how to implement the methodology that we propose. First, we state the likelihood implied by the VAR given in equation (1). Second, we state the conjugate prior and the posterior in the unrestricted version of that VAR. Third, we describe how to evaluate the marginal likelihood implied by the VAR given in equation (1) without restrictions, with a single block-exogeneity restriction, and with multiple block-exogeneity restrictions. Fourth, we discuss how to search for the best VAR in the family Ω when the family Ω is too large to evaluate the marginal likelihood implied by each VAR in that family.

3.1 Likelihood

The likelihood of the data implied by the VAR given in equation (1), conditional on initial observations, is

$$p(Y|B, \Sigma) = (2\pi)^{-NT/2} |\Sigma|^{-T/2} \exp \left(-\frac{1}{2} \text{tr}(Y - XB)'(Y - XB)\Sigma^{-1} \right), \quad (5)$$

where N is the length of the vector $y(t)$, T is the number of observations in the sample,

$$Y_{T \times N} = \begin{pmatrix} y(1)' \\ y(2)' \\ \vdots \\ y(T)' \end{pmatrix}, \quad B_{K \times N} = \begin{pmatrix} B_1' \\ \vdots \\ B_P' \\ \gamma' \end{pmatrix},$$

$K = NP + 1$, and

$$X_{T \times K} = \begin{pmatrix} y(0)' & y(-1)' & \dots & y(1-P)' & 1 \\ y(1)' & y(0)' & \dots & y(2-P)' & 1 \\ \vdots & \vdots & & \vdots & \vdots \\ y(T-1)' & y(T-2)' & \dots & y(T-P)' & 1 \end{pmatrix}.$$

3.2 Conjugate prior and posterior in the unrestricted model

Consider the prior density of B and Σ in the unrestricted model $\omega^U \in \Omega$. We assume that $p(B, \Sigma|\omega^U)$ is conjugate, that is

$$p(B, \Sigma|\omega^U) \propto |\Sigma|^{-(\tilde{\nu} + K + N + 1)/2} \exp \left(-\frac{1}{2} \text{tr}(\tilde{Y} - \tilde{X}B)'(\tilde{Y} - \tilde{X}B)\Sigma^{-1} \right), \quad (6)$$

where $\tilde{\nu}$, \tilde{Y} , and \tilde{X} are prior hyperparameters of appropriate dimensions.¹¹ Let

$$\tilde{Q} = (\tilde{X}'\tilde{X})^{-1}, \quad \tilde{B} = (\tilde{X}'\tilde{X})^{-1}\tilde{X}'\tilde{Y}, \quad \text{and} \quad \tilde{S} = (\tilde{Y} - \tilde{X}\tilde{B})'(\tilde{Y} - \tilde{X}\tilde{B}).$$

It is straightforward to show that, so long as $\tilde{\nu} > 0$, this prior is proper and satisfies

$$p(B, \Sigma | \omega^U) = p(B | \Sigma, \omega^U) p(\Sigma | \omega^U) = \mathcal{N}(\text{vec } \tilde{B}, \Sigma \otimes \tilde{Q}) \mathcal{IW}(\tilde{S}, \tilde{\nu}), \quad (7)$$

where \mathcal{N} denotes a multivariate normal density and \mathcal{IW} denotes an inverted Wishart density.¹²

When we combine prior (6) with likelihood (5), we obtain the posterior of B and Σ in the unrestricted model in the family Ω . Let $\bar{\nu} = \tilde{\nu} + T$,

$$\bar{Y} = \begin{pmatrix} \tilde{Y} \\ Y \end{pmatrix} \quad \bar{X} = \begin{pmatrix} \tilde{X} \\ X \end{pmatrix},$$

$$\bar{Q} = (\bar{X}'\bar{X})^{-1}, \quad \bar{B} = (\bar{X}'\bar{X})^{-1}\bar{X}'\bar{Y}, \quad \text{and} \quad \bar{S} = (\bar{Y} - \bar{X}\bar{B})'(\bar{Y} - \bar{X}\bar{B}).$$

It is straightforward to show that, so long as $\bar{\nu} > 0$, the posterior is proper and satisfies

$$p(B, \Sigma | Y, \omega^U) = p(B | \Sigma, Y, \omega^U) p(\Sigma | Y, \omega^U) = \mathcal{N}(\text{vec } \bar{B}, \Sigma \otimes \bar{Q}) \mathcal{IW}(\bar{S}, \bar{\nu}). \quad (8)$$

3.3 Computation of marginal likelihood

We now describe how to evaluate the marginal likelihood of the data implied by the VAR given in equation (1) without restrictions, with a single block-exogeneity restriction, and with multiple block-exogeneity restrictions. In other words, we describe how to evaluate the prior predictive density of the data implied by each model $\omega \in \Omega$.

We adopt the following approach. First, we evaluate the marginal likelihood implied by the unrestricted model ω^U , $p(Y | \omega^U)$. Second, for a given restricted model ω^R , we evaluate the Bayes factor in favor of that model relative to the unrestricted model. This Bayes factor is equal to $p(Y | \omega^R) / p(Y | \omega^U)$. Third, we multiply this Bayes factor by $p(Y | \omega^U)$ to obtain $p(Y | \omega^R)$, the marginal likelihood implied by the restricted model ω^R .

¹¹Section 5.2 and Appendix A discuss the specification of the prior hyperparameters in this paper's application.

¹²See Bauwens et al. (1999), Appendix A, for the definitions of the multivariate normal and inverted Wishart densities.

It is well known how to evaluate the marginal likelihood of the data implied by an unrestricted VAR with the conjugate prior. Namely, the marginal likelihood of the data implied by the model $\omega^U \in \Omega$ can be evaluated exactly based on the following expression

$$p(Y|\omega^U) = \pi^{-NT/2} \frac{\Gamma_N\left(\frac{\bar{\nu}}{2}\right) |\tilde{X}'\tilde{X}|^{N/2} |\tilde{S}|^{\bar{\nu}/2}}{\Gamma_N\left(\frac{\bar{\nu}}{2}\right) |\tilde{X}'\tilde{X}|^{N/2} |\tilde{S}|^{\bar{\nu}/2}}, \quad (9)$$

where Γ_N denotes the multivariate Gamma function.¹³

Next, we show that it is straightforward to evaluate the Bayes factors in favor of models with block-exogeneity restrictions relative to the unrestricted model.

Savage-Dickey result. We observe that one can use the Savage-Dickey result of Dickey (1971) to evaluate the Bayes factor in favor of a VAR with block-exogeneity relative to an unrestricted VAR.

Consider a restricted model $\omega^R \in \Omega$. The Savage-Dickey result states that if

$$p(B_U, \Sigma|\omega^R) = p(B_U, \Sigma|\omega^U, B_R = 0), \quad (10)$$

the Bayes factor in favor of the restricted model has the property

$$\frac{p(Y|\omega^R)}{p(Y|\omega^U)} = \frac{p(B_R = 0|Y, \omega^U)}{p(B_R = 0|\omega^U)}. \quad (11)$$

Here $p(B_R = 0|Y, \omega^U)$ is the marginal posterior density of B_R in the unrestricted model, evaluated at the point $B_R = 0$. Furthermore, $p(B_R = 0|\omega^U)$ is the marginal prior density of B_R in the unrestricted model, evaluated at the point $B_R = 0$. In words, the Savage-Dickey result states that the Bayes factor for the test of the restriction $B_R = 0$ against the alternative $B_R \neq 0$ is equal to the ratio of the marginal posterior density of B_R evaluated at zero to the marginal prior density of B_R evaluated at zero. The right-hand-side of expression (11) is known as the Savage-Dickey ratio. Recall that the prior density $p(B_U, \Sigma|\omega^R)$ in any restricted model $\omega^R \in \Omega$ has property (3), which is the same as condition (10). Therefore, the Savage-Dickey result applies to tests of block-exogeneity restrictions in the family Ω .

Single block-exogeneity restriction. When we test a single block-exogeneity restriction, the Savage-Dickey ratio is available analytically.

¹³See Technical Appendix for a derivation of equation (9) and the definition of the multivariate Gamma function.

Consider a partition of the vector y modeled in equation (1) into two subvectors, $y = \{y_i, y_j\}$. Suppose that we want to test if y_i is block-exogenous to y_j . Let $\omega' \in \Omega$ be the model in which this block-exogeneity restriction holds and there are no other restrictions. Let α denote the column indices of the variables represented by y_i in matrix Y . Let β denote the column indices of the lags of the variables represented by y_j in matrix X . The block-exogeneity restriction is $B_{\beta,\alpha} = 0$.¹⁴ The marginal prior density of $B_{\beta,\alpha}$ in the unrestricted model is

$$p(B_{\beta,\alpha}|\omega^U) = \mathcal{T}(\tilde{B}_{\beta,\alpha}, (\tilde{Q}_{\beta,\beta})^{-1}, \tilde{S}_{\alpha,\alpha}, \tilde{\nu} - N_j), \quad (12)$$

where \mathcal{T} denotes a matrixvariate Student density and N_j is the number of variables in y_j .¹⁵ Furthermore, the marginal posterior density of $B_{\beta,\alpha}$ in the unrestricted model is

$$p(B_{\beta,\alpha}|Y, \omega^U) = \mathcal{T}(\bar{B}_{\beta,\alpha}, (\bar{Q}_{\beta,\beta})^{-1}, \bar{S}_{\alpha,\alpha}, \bar{\nu} - N_j). \quad (13)$$

The analytical expressions (12)-(13) are available due to the fact that when we pick an intersection of rows β and columns α from B , the variance-covariance matrix of the resulting vector $\text{vec } B_{\beta,\alpha}$ has the Kronecker structure just like the variance-covariance matrix of $\text{vec } B$.

Using the definition of the matrixvariate Student density, evaluating densities (12) and (13) at the point $B_{\beta,\alpha} = 0$ and substituting into equality (11), we find that the Bayes factor for the test of the restriction $B_{\beta,\alpha} = 0$ against the alternative $B_{\beta,\alpha} \neq 0$ is equal to

$$\begin{aligned} \frac{p(Y|\omega')}{p(Y|\omega^U)} &= \frac{\Gamma_{N_i} \left(\frac{\bar{\nu} - N_j + K_j}{2} \right)}{\Gamma_{N_i} \left(\frac{\bar{\nu} - N_j}{2} \right)} \frac{\Gamma_{N_i} \left(\frac{\tilde{\nu} - N_j}{2} \right)}{\Gamma_{N_i} \left(\frac{\tilde{\nu} - N_j + K_j}{2} \right)} \\ &\times \frac{|\bar{S}_{\alpha,\alpha}|^{\frac{\bar{\nu} - N_j}{2}} |(\bar{Q}_{\beta,\beta})^{-1}|^{\frac{N_i}{2}} |\bar{S}_{\alpha,\alpha} + \bar{B}'_{\beta,\alpha} (\bar{Q}_{\beta,\beta})^{-1} \bar{B}_{\beta,\alpha}|^{-\frac{\bar{\nu} - N_j + K_j}{2}}}{|\tilde{S}_{\alpha,\alpha}|^{\frac{\tilde{\nu} - N_j}{2}} |(\tilde{Q}_{\beta,\beta})^{-1}|^{\frac{N_i}{2}} |\tilde{S}_{\alpha,\alpha} + \tilde{B}'_{\beta,\alpha} (\tilde{Q}_{\beta,\beta})^{-1} \tilde{B}_{\beta,\alpha}|^{-\frac{\tilde{\nu} - N_j + K_j}{2}}}. \end{aligned} \quad (14)$$

In expression (14), N_i denotes the number of variables in the vector y_i and K_j is equal to the product of the number of variables in the vector y_j and the number of lags in the model.¹⁶

¹⁴ $B_{\beta,\alpha}$ is a matrix consisting of the intersection of rows β and columns α of the matrix B .

¹⁵ See Bauwens et al. (1999), Appendix A.2.7, for the definition of the matrixvariate Student density and a proof of equality (12).

¹⁶ $B_{\beta,\alpha}$ has size $K_j \times N_i$.

Multiple block-exogeneity restrictions. When we test multiple block-exogeneity restrictions simultaneously, we can approximate the Savage-Dickey ratio numerically with a simple Monte Carlo.

We focus here on the case of two block-exogeneity restrictions. A generalization is straightforward. Consider a partition of the vector y modeled in equation (1) into three subvectors, $y = \{y_i, y_j, y_k\}$. Suppose that we want to test if y_i is block-exogenous to y_j and y_k , and y_i and y_j are block-exogenous to y_k . Let α_1 denote the column indices of the variables represented by y_i in matrix Y . Let β_1 denote the column indices of the lags of the variables represented by y_j and y_k in matrix X . Let α_2 denote the column indices of the variables represented by y_j in matrix Y . Let β_2 denote the column indices of the lags of the variables represented by y_k in matrix X . The block-exogeneity restrictions are $B_{\beta_1, \alpha_1} = 0$ and $B_{\beta_2, \alpha_2} = 0$.

The marginal density of $\text{vec } B_{\beta_1, \alpha_1}$ and the marginal density of $\text{vec } B_{\beta_2, \alpha_2}$ are each multivariate Student, but the joint density of $((\text{vec } B_{\beta_1, \alpha_1})', (\text{vec } B_{\beta_2, \alpha_2})')'$, which enters the Savage-Dickey ratio, is not available analytically. In particular, this joint density is not multivariate Student because the variance-covariance matrix of $((\text{vec } B_{\beta_1, \alpha_1})', (\text{vec } B_{\beta_2, \alpha_2})')'$ does not have the Kronecker structure.¹⁷

However, the density of $((\text{vec } B_{\beta_1, \alpha_1})', (\text{vec } B_{\beta_2, \alpha_2})')'$ conditional on Σ is multivariate normal. The reason is that $((\text{vec } B_{\beta_1, \alpha_1})', (\text{vec } B_{\beta_2, \alpha_2})')'$ is a subvector of $\text{vec } B$ which, conditionally on Σ , is multivariate normal. See equations (7) and (8). Therefore, the marginal prior density at zero can be approximated from M Monte Carlo draws of Σ from its prior $p(\Sigma|\omega^U)$ as

$$\begin{aligned} & p(((\text{vec } B_{\beta_1, \alpha_1})', (\text{vec } B_{\beta_2, \alpha_2})')' = 0 | \omega^U) \\ &= \frac{1}{M} \sum_{m=1}^M p(((\text{vec } B_{\beta_1, \alpha_1})', (\text{vec } B_{\beta_2, \alpha_2})')' = 0 | \Sigma^m, \omega^U). \end{aligned}$$

Analogously, the marginal posterior density at zero can be approximated from M Monte Carlo draws of Σ from its posterior $p(\Sigma|Y, \omega^U)$ as

$$p(((\text{vec } B_{\beta_1, \alpha_1})', (\text{vec } B_{\beta_2, \alpha_2})')' = 0 | Y, \omega^U)$$

¹⁷The joint density of $((\text{vec } B_{\beta_1, \alpha_1})', (\text{vec } B_{\beta_2, \alpha_2})')'$ is not multivariate Student, either. See Appendix A of Bauwens et al. (1999) for the definitions and properties of matrix-variate and multivariate Student densities.

$$= \frac{1}{M} \sum_{m=1}^M p \left(((\text{vec } B_{\beta_1, \alpha_1})', (\text{vec } B_{\beta_2, \alpha_2})')' = 0 | \Sigma^m, Y, \omega^U \right).$$

3.4 Finding the best model when the family Ω is too large to check all models

In principle, one can evaluate the marginal likelihood of the data implied by each VAR in the family Ω proceeding as in Section 3.3. However, the number of elements in Ω grows very quickly with N_2 . Consider the number of block-exogeneity restrictions among N_2 variables, denoted $C(N_2)$. Since block-exogeneity is a transitive relation, $C(N_2)$ is equal to the number of weak orders of N_2 elements. One can show¹⁸ that

$$C(N_2) = \sum_{k=0}^{N_2-1} \binom{N_2}{k} C(k),$$

and for a large N_2

$$C(N_2) \approx \frac{N_2!}{2(\ln 2)^{N_2+1}}.$$

The number of elements in Ω is $K(N_2) = 2C(N_2)$, that is, the number of elements in Ω is equal to twice the number of block-exogeneity restrictions among the N_2 variables represented by y_2 .¹⁹ For example: $K(N_2 = 2) = 6$, $K(N_2 = 3) = 26$, $K(N_2 = 4) = 150$, $K(N_2 = 5) = 1082$, $K(N_2 = 6) = 9366$, $K(N_2 = 7) = 94586$, and so on.

In practice, $K(N_2)$ quickly becomes too large for us to compute the marginal likelihood implied by each model in the family Ω – our computers are too slow. Therefore, we implement the Markov chain Monte Carlo model composition (MC³) algorithm of Madigan and York (1995) to select a subset of models from Ω , and we only compute the marginal likelihood implied by each model in this subset. We use the MC³ algorithm because this algorithm samples models from their posterior distribution. In particular, the MC³ algorithm samples models with higher posterior probabilities more frequently than models with lower posterior probabilities. Therefore, when we use this algorithm, we spend most time evaluating marginal likelihood for well-fitting models and we spend little time evaluating marginal likelihood for models with poor fit. The details are in Appendix B.

¹⁸See OEIS (2011).

¹⁹The reason why multiplication by two is necessary is that, given each pattern of block-exogeneity restrictions within y_2 , we can either have block-exogeneity between y_1 and y_2 or not.

4 Alternative approaches

In this section, we compare our methodology with three alternative methodologies to answer the question that we study in this paper. We argue that the alternative methodologies are less attractive than the methodology that we propose.

The question that we study in this paper is a question concerning model choice. It is familiar that a Bayesian answer to a question concerning model choice rests on computation of marginal likelihood – prior predictive density – of the data. The challenge in the question that we study is twofold. One must specify: (i) the prior predictive density of *which data* is to be computed, and (ii) how to compute this prior predictive density. Note that it does not make sense to compare the marginal likelihood of $(Y_1, Y_{2,i})$ implied by a model of y_1 and a subvector of y_2 called $y_{2,i}$ with the marginal likelihood of $(Y_1, Y_{2,j})$ implied by a model of y_1 and another subvector of y_2 called $y_{2,j}$. This would be like comparing apples with oranges. We observe that evaluating certain restrictions in the joint model of y_1 and y_2 amounts to evaluating whether subvectors of y_2 are helpful in modeling y_1 . The practical implication is that we always compute *the marginal likelihood of $Y = (Y_1, Y_2)$, without restrictions or with restrictions*. In Section 3, we explained how to compute this marginal likelihood. We now discuss three alternative methodologies such that each methodology involves computation of *a predictive density of Y_1 only*.

Let ψ denote a model of y_1 and a subvector of y_2 called $y_{2,\psi}$. Consider the following three statistics. The first statistic is the marginal predictive density of Y_1 , that is, the marginal likelihood of $(Y_1, Y_{2,\psi})$ marginalized with respect to $Y_{2,\psi}$

$$p(Y_1|\psi) = \int p(Y_1, Y_{2,\psi}|\psi) dY_{2,\psi}. \quad (15)$$

The second statistic is the predictive density of Y_1 conditional on the actually observed $Y_{2,\psi}$

$$p(Y_1|Y_{2,\psi}, \psi) = \frac{p(Y_1, Y_{2,\psi}|\psi)}{\int p(Y_1, Y_{2,\psi}|\psi) dY_1}. \quad (16)$$

The third statistic is the predictive density score²⁰ of Y_1 at horizon $h > 0$, typically computed as

$$g(Y_1, h|\psi) = \prod_{t=1}^{T-h} p(y_1(t+h)|y(\tau : \tau \leq t), \psi). \quad (17)$$

²⁰The predictive density score is used in many papers. See Geweke and Amisano (2011) for a discussion.

Let us understand what each of these statistics tells us and compare these statistics to the statistic that we use. We first present four useful expressions, and then we discuss these expressions. For the marginal predictive density of Y_1 , we have

$$\begin{aligned} p(Y_1|\psi) &= p(y_1(1, \dots, T)|y_1(-P+1, \dots, 0), y_{2,\psi}(-P+1, \dots, 0), \psi) \\ &= \prod_{j=1}^Q p(y_1(s_{j-1}+1, \dots, s_j)|y_1(-P+1, \dots, s_{j-1}), y_{2,\psi}(-P+1, \dots, 0), \psi). \end{aligned} \quad (18)$$

Here we partition the sequence of dates $0, 1, \dots, T$ using a strictly increasing sequence of integers $\{s_j\}_{j=0}^Q$ with $s_0 = 0$ and $s_Q = T$.²¹ Furthermore, we make explicit the conditioning on the P initial observations $y_1(-P+1, \dots, 0)$ and $y_{2,\psi}(-P+1, \dots, 0)$. For the predictive density of Y_1 conditional on the actually observed $Y_{2,\psi}$, we have

$$\begin{aligned} p(Y_1|Y_2, \psi) &= p(y_1(1, \dots, T)|y_1(-P+1, \dots, 0), y_{2,\psi}(-P+1, \dots, T), \psi) \\ &= \prod_{j=1}^Q p(y_1(s_{j-1}+1, \dots, s_j)|y_1(-P+1, \dots, s_{j-1}), y_{2,\psi}(-P+1, \dots, T), \psi). \end{aligned} \quad (19)$$

For the predictive density score of Y_1 , we have²²

$$g(Y_1, \{s_j\}_{j=0}^Q|\psi) = \prod_{j=1}^Q p(y_1(s_{j-1}+1, \dots, s_j)|y_1(-P+1, \dots, s_{j-1}), y_{2,\psi}(-P+1, \dots, s_{j-1}), \psi). \quad (20)$$

For the statistic that we use, the marginal likelihood of Y implied by a model in the family Ω , we have

$$\begin{aligned} p(Y|\omega) &= p(y_1(1, \dots, T), y_2(1, \dots, T)|y_1(-P+1, \dots, 0), y_2(-P+1, \dots, 0), \omega) \\ &= \prod_{j=1}^Q p(y_1(s_{j-1}+1, \dots, s_j), y_2(s_{j-1}+1, \dots, s_j)|y_1(-P+1, \dots, s_{j-1}), y_2(-P+1, \dots, s_{j-1}), \omega) \end{aligned} \quad (21)$$

The following lessons emerge from comparing equations (18)-(20) with equation (21).

The marginal predictive density of Y_1 , $p(Y_1|\psi)$, measures the out-of-sample fit to the data on y_1 *assuming that no data on $y_{2,\psi}$ are available except for the initial observations.*

²¹This partitioning follows Geweke (2005), p.67.

²²Expression (20) is a valid way to define the predictive density score, alternative to (17). Expressions (17) and (20) coincide when $h = 1$ in (17) and $\{s_j\}_{j=0}^Q = \{0, 1, \dots, T\}$ in (20). We think that expression (20) makes more transparent the comparison between the predictive density score and the other statistics we consider here.

Note the term $y_{2,\psi}(-P+1, \dots, 0)$ in expression (18). The fact that this statistic discards all available data on $y_{2,\psi}$ except for the initial observations makes this statistic unattractive as a criterion for model choice. Consider the following, fairly common case. Suppose that we want to compare a VAR model ψ of y_1 and $y_{2,\psi}$ with another VAR model ψ' of y_1 and another subvector of y_2 called $y_{2,\psi'}$, where $y_{2,\psi'}$ has the same number of variables as $y_{2,\psi}$. Each VAR has one lag and the same prior, e.g. the standard Sims-Zha prior.²³ Suppose that we rescale variables such that each variable in $y_{2,\psi}$ and each variable in $y_{2,\psi'}$ have the same value in period $t = 0$.²⁴ Then it is straightforward to show that $p(Y_1|\psi) = p(Y_1|\psi')$, that is, the marginal predictive density of Y_1 implied by the model ψ is equal to the marginal predictive density of Y_1 implied by the model ψ' . The implication is strong. If we used $p(Y_1|\psi)$ to decide whether to include $y_{2,\psi}$ or $y_{2,\psi'}$ in the VAR with y_1 , we would end up indifferent. Even if $y_{2,\psi}$ were strongly related to y_1 and $y_{2,\psi'}$ followed a white noise process.²⁵

The predictive density of Y_1 conditional on the actually observed $Y_{2,\psi}$, $p(Y_1|Y_{2,\psi}, \psi)$, measures the fit to the data on y_1 *assuming that data on $y_{2,\psi}$ have been observed through the end of the sample, period T* . Note the term $y_{2,\psi}(-P+1, \dots, T)$ in expression (19). Thus $p(Y_1|Y_{2,\psi}, \psi)$ is not an out-of-sample measure of fit. This statistic tells us only how well a model ψ captures the relation between $y_{2,\psi}$ and y_1 , for a particular $Y_{2,\psi}$, namely the actually observed $Y_{2,\psi}$. A model ψ could attain a high value of $p(Y_1|Y_{2,\psi}, \psi)$ while yielding poor out-of-sample fit to the data on $y_{2,\psi}$ and, therefore, also poor out-of-sample fit to the data on y_1 . This feature makes $p(Y_1|Y_{2,\psi}, \psi)$ unattractive as a criterion for model choice.

The predictive density score of Y_1 , $g(Y_1, \{s_j\}_{j=0}^Q|\psi)$, measures the out-of-sample fit to the data on y_1 . The predictive prior density of Y , $p(Y|\omega)$, measures the out-of-sample fit to the data on y . Both statistics condition on all data available at the time when a density

²³Concerning the Sims-Zha prior see Section 5.2 and Appendix A.

²⁴For example, suppose that $y_{2,\psi}$ consists of a single variable, $y_{2,\psi'}$ consists of a single variable, $y_{2,\psi}(0) = 5$, and $y_{2,\psi'}(0) = 10$. Then multiplication of $y_{2,\psi}$ by 2 yields $y_{2,\psi}(0) = y_{2,\psi'}(0) = 10$, that is, $y_{2,\psi}$ and $y_{2,\psi'}$ have the same value in period $t = 0$.

²⁵If one used a training sample prior in addition to the standard Sims-Zha prior, the marginal predictive densities of Y_1 in this example would not be exactly equal to each other. In our application, we computed the marginal predictive density of Y_1 implied by many VARs, with a training sample prior in addition to the standard Sims-Zha prior. We found that the differences between the marginal predictive densities of Y_1 across different VARs were very small.

is being evaluated. Note the term $y_1(-P+1, \dots, s_{j-1}), y_{2,\psi}(-P+1, \dots, s_{j-1})$ in expression (20) and the term $y_1(-P+1, \dots, s_{j-1}), y_2(-P+1, \dots, s_{j-1})$ in expression (21). This feature is attractive and distinguishes both statistics from $p(Y_1|\psi)$ and $p(Y_1|Y_{2,\psi}, \psi)$.

The following three differences between $g(Y_1, \{s_j\}_{j=0}^Q|\psi)$ and $p(Y|\omega)$ make the former statistic unattractive as a criterion for model choice.

First, computation of $p(Y|\omega)$ leads directly to computation of posterior odds on models. See Sections 2 and 3. In contrast, one cannot assign probabilities to models based on $g(Y_1, \{s_j\}_{j=0}^Q|\psi)$.²⁶

Second, any partition $\{s_j\}_{j=0}^Q$ leads to the same value of $p(Y|\omega)$ for a given model. In contrast, different partitions $\{s_j\}_{j=0}^Q$ lead in general to different values of $g(Y_1, \{s_j\}_{j=0}^Q|\psi)$ for a given model. For example, consider the partition $\{0, 1, 2, \dots, T\}$ that decomposes $p(Y|\omega)$ and $g(Y_1, \{s_j\}_{j=0}^Q|\psi)$ into one-step-ahead predictive densities; and consider the partition $\{0, 4, 8, \dots, T\}$ that decomposes $p(Y|\omega)$ and $g(Y_1, \{s_j\}_{j=0}^Q|\psi)$ into one-to-four-steps-ahead predictive densities. Both partitions yield the same value of $p(Y|\omega)$ but different values of $g(Y_1, \{s_j\}_{j=0}^Q|\psi)$. The practical implication is that according to $g(Y_1, \{s_j\}_{j=0}^Q|\psi)$ the best model for forecasting one period ahead in general differs from the best model for forecasting one-to-four periods ahead. Thus model choice with this statistic requires an arbitrary weighting of forecast horizons.^{27,28}

Third, computation of $g(Y_1, \{s_j\}_{j=0}^Q|\psi)$ requires looping over $\{s_j\}_{j=0}^Q$ and, for each s_j , evaluating a predictive density of $y_1(s_{j-1} + 1, \dots, s_j)$. This computation can be time consuming, in particular when T is large or $s_j - s_{j-1}$ is large. In contrast, in the case of a single block-exogeneity restriction $p(Y|\omega)$ is available analytically and in the case of multiple block-exogeneity restrictions only a simple Monte Carlo is required to evaluate $p(Y|\omega)$.

²⁶One can assign probabilities to models based on $p(Y_1|\psi)$. These probabilities are conditional on the actually observed Y_1 . In contrast, $p(Y|\omega)$ yields probabilities conditional on a larger information set, the actually observed Y . One cannot assign probabilities to models based on $p(Y_1|Y_{2,\psi}, \psi)$. It turns out that $p(Y_1|Y_{2,\psi}, \psi)$ is proportional to the change in probability of model ψ once Y_1 has been observed in addition to $Y_{2,\psi}$.

²⁷In practice, researchers compute predictive density scores not only for partitions like $\{s_j\}_{j=0}^Q$, but also for different horizons, like in equation (17) when $h > 1$. Each horizon h leads to a different value of $g(Y_1, h | \psi)$.

²⁸Any partition $\{s_j\}_{j=0}^Q$ leads to the same value of $p(Y_1|\psi)$ and the same value of $p(Y_1|Y_{2,\psi}, \psi)$ for a given model.

See Section 3.3. The computational burden of this Monte Carlo increases with N whereas the computational burden of evaluating $g(Y_1, \{s_j\}_{j=0}^Q | \psi)$ increases with T . Therefore, in large samples computation of $p(Y|\omega)$ is guaranteed to be cheaper than computation of $g(Y_1, \{s_j\}_{j=0}^Q | \psi)$.²⁹

5 Application

We turn to an application of our methodology. In this section we study the following question. We want to estimate a VAR with a measure of the price level, GDP, and a short-term interest rate in the euro area. Which other macroeconomic and financial variables, if any, should we include in that VAR?

We are interested in this question, because we seek a benchmark model with the following variables of interest: the price level, GDP, and an interest rate controlled by monetary policy in the euro area. We think of the European Central Bank (ECB) as affecting the price level via changes in a policy-controlled interest rate. Therefore, the benchmark model must include a measure of the price level and an interest rate controlled by the ECB. We also think that when the ECB evaluates different paths of the policy rate, the ECB considers implications for real economic activity. Therefore, we include GDP as a measure of real economic activity in the benchmark model. We want the benchmark model to fit these three variables well out-of-sample; and we want a methodology that lets us establish which other variables, if any, improve the out-of-sample fit to these three variables. We seek a benchmark model, rather than average over many models, because it is much easier to maintain and communicate to policymakers results from a benchmark model. Furthermore, in the future we aim to build on this paper and study the effects of monetary policy shocks on the price level and GDP. We want identification of monetary policy to be based on stochastic prior restrictions. For specification of the prior and for computational feasibility, it will be important to use a single model, as opposed to averaging over many models, it will be important that the model be of medium-size (that is, not too large), and it will be

²⁹Computation of $p(Y_1|\psi)$ and $p(Y_1|Y_{2,\psi}, \psi)$ requires evaluating the marginal likelihood of an unobserved components model. This computation is difficult when there are many unobservable state variables. Furthermore, the computational burden of evaluating $p(Y_1|\psi)$ and $p(Y_1|Y_{2,\psi}, \psi)$ increases with T .

important that the model be a VAR (that is, not a complex non-linear model).

5.1 Dataset

We put together a dataset with $N = 20$ variables, of which $N_1 = 3$ are the variables of interest and $N_2 = 17$ are the remaining variables. Table 1 lists the variables in our dataset, with units of measurement and any transformations.³⁰

The variables of interest, numbered 1-3 in Table 1, are the Harmonized Index of Consumer Prices (HICP), real GDP, and the overnight interbank interest rate Eonia. We think of the ECB as controlling the Eonia. Out of many possible remaining variables, in this paper we focus on the following variables.³¹ All variables are for the euro area as a whole, except when indicated. We include two components of GDP: real consumption and real investment. We include the unemployment rate as a measure of capacity utilization, a notion central to Keynesian and New Keynesian business cycle models. We include the yield on 10-year government bonds as a measure of the long-term interest rate (Bond Yield 10y). We include the spread between corporate bonds rated BBB with maturity 7-10 years and government bonds with the same maturity, a measure of the credit spread central to business cycle models with financial frictions (Bond Spread BBB 7to10y). We include a measure of money supply (M3), money being a variable central to monetarist business cycle models. We include loans to non-financial corporations as a measure of the quantity of credit (Loans NFC). We include the nominal effective exchange rate of the euro (Effective Exchange Rate). We include a measure of the world price of oil (Oil Prices) and an index of world commodity prices (Commodity Prices), both measured in U.S. dollars. We include two measures of activity in the housing market: an index of house prices (House Prices) and real housing investment. We include two measures of activity in stock markets: the Dow Jones EuroStoxx index and the VStoxx implied volatility index. We include the most

³⁰The source of the dataset is the database of the ECB. The dataset is available from the authors upon request.

³¹In essentially any application, the choice of the dataset will be informal, based on the researcher's prior knowledge, because the set of all possible remaining variables is extremely large and most of those variables can be seen to be a priori irrelevant for the question of interest. The methodology that we propose formalizes the choice of variables to include in a model with variables of interest once a particular dataset is available.

	Variable	Units	Trsf	
1	HICP	index	SA	log
2	Real GDP	2000 Euro millions	SA	log
3	Eonia	percent p.a.		none
4	Real Consumption	2000 Euro millions	SA	log
5	Real Investment	2000 Euro millions	SA	log
6	Unemployment Rate	percent of civilian workforce	SA	none
7	Bond Yield 10y	percent p.a.		none
8	Bond Spread BBB 7to10y	percentage points		none
9	M3	Euro millions	SA	log
10	Loans NFC	Euro millions		log
11	Effective Exchange Rate	index		log
12	Oil Price	US dollar per barrel		log
13	Commodity Prices	index		log
14	House Prices	index	SA	log
15	Real Housing Investment	2000 Euro millions	SA	log
16	Euro Stoxx	index		log
17	VStoxx	percent p.a.		log
18	PMI	index, 50 = no change		log
19	US Real GDP	2000 dollar billions	SA	log
20	US Fed Funds Rate	percent		none

Table 1: Variable names, units and transformations.

popular leading indicator of economic activity, the purchasing managers index (PMI). Finally, we include two U.S. variables: GDP and the federal funds rate, as measures of the effect of the U.S. economy on the euro area.

5.2 Prior

We construct our prior in two steps: (i) we start with an initial prior formulated before seeing any data, and (ii) we combine the initial prior with a training sample prior. Formally, matrices \tilde{Y}, \tilde{X} in expression (6) consist of two blocks

$$\tilde{Y} = \begin{pmatrix} Y_{SZ} \\ Y_{ts} \end{pmatrix}, \quad \tilde{X} = \begin{pmatrix} X_{SZ} \\ X_{ts} \end{pmatrix}, \quad (22)$$

where the terms Y_{SZ} , Y_{ts} , X_{SZ} , and X_{ts} are defined below.

The initial prior is the prior proposed by Sims and Zha (1998). We implement the Sims-Zha prior by creating dummy observations Y_{SZ} and X_{SZ} . The term $\tilde{\nu}$ in expression (6) also belongs to the initial prior. The Sims-Zha prior is controlled by several hyperparameters. We set the following values for the key hyperparameters: the “overall tightness” is set to 0.1, the weight of the “one-unit-root” dummy is set to 1, and the weight of the “no-cointegration dummy” is set to 1. Appendix A gives the details concerning the Sims-Zha prior and explains our choice of hyperparameter values.

Our sample contains quarterly data from 1999Q1 to 2010Q4. We also use a training sample. In addition to the Sims-Zha prior, we add to the prior the information from the pre-EMU period 1989Q1 to 1998Q4. Y_{ts} and X_{ts} denote the matrices with the data from this training sample. We found that adding this training sample improves the marginal likelihood in the sample 1999Q1-2010Q4 compared with using the Sims-Zha prior only.

5.3 Findings

The VAR that we focus on has one lag, that is, $P = 1$. We found that including more lags reduces the marginal likelihood.

Main finding: which variables to include in the VAR? By assumption, the first block of the best model in the family Ω , ω^* , includes the price level, GDP, and the policy rate. We find that the first block of the best model ω^* in addition includes the following four

rank	variables from y_2 in the first block	odds to ω^*
1 (ω^*)	6,8,18,20	1.00
2	6,8,18,20	0.50
3	6,8,18,20	0.37
4	6,8,18,20	0.37
5	6,8,18,20	0.30
6	6,7,8,18,20	0.30
7	6,7,8,18,20	0.27
8	6,7,8,18,20	0.27
9	6,8,18,20	0.27
10	6,7,8,18,20	0.25

Table 2: Best ten models.

variables: the unemployment rate, the bond spread, the purchasing managers index, and the federal funds rate. Thus one main finding is the following. If one wants to model the price level, GDP, and the policy rate in the euro area in a VAR, that VAR should also include a measure of capacity utilization, a notion central to Keynesian and New Keynesian business cycle models (the unemployment rate), a variable central to business cycle models with financial frictions (the bond spread), the main leading indicator (the purchasing managers index), and a variable external to the euro area (the federal funds rate). We think that this is a plausible finding.

Table 2 reports the first block of each of the best ten models in the family Ω . For each model, the table reports the posterior odds in favor of that model relative to the best model ω^* . The best model ω^* is in the first row of the table.³²

Consider Table 2. The top five models have the same remaining variables in the first block: the unemployment rate, the bond spread, the purchasing managers index, and the

³²Table 2 and the next table use variable numbers from Table 1 instead of variable names, in order to conserve space. For example, “6, 8, 18, 20” next to the model ω^* stands for “Unemployment Rate, Bond Spread BBB 7to10y, PMI, and US Fed Funds Rate”.

federal funds rate. Furthermore, six out of the top ten models have *only* those same remaining variables in the first block. The other four out of the top ten models also have those same remaining variables in the first block *plus only one other variable*, the long-term interest rate.

The best model ω^* has multiple block-exogeneity restrictions. In addition to the block-exogeneity restriction between the seven variables in the first block and all the other variables, the model ω^* has seven additional block-exogeneity restrictions. Thus the model ω^* consists of nine blocks. Four of these blocks consist of two variables and another four of these blocks consist of one variable.³³

The best model versus the unrestricted model. The data strongly support the block-exogeneity restrictions in the best model ω^* . The posterior odds in favor of the model ω^* relative to the unrestricted model ω^U are approximately 5×10^7 . If we think of the first block of the model ω^* as a medium-size VAR and we think of the unrestricted model as a large VAR, we find strong support for the medium-size VAR.³⁴

Tightening the Sims-Zha prior is no substitute for the block-exogeneity restrictions. We consider the model with the same block-exogeneity restrictions as the best model ω^* and we tighten the Sims-Zha prior by reducing the “overall tightness” hyperparameter gradually from 0.1 (the best model) to 0.005. The posterior odds relative to the unrestricted model fall but remain enormous, about 1×10^6 . Furthermore, tightening the Sims-Zha prior produces a much worse model: the posterior odds in favor of the model ω^* (with the “overall tightness” hyperparameter set to 0.1) relative to the model with the same block-exogeneity restrictions and the “overall tightness” hyperparameter set to 0.005 are enormous, about 6×10^{18} .³⁵

³³The second block of the model ω^* consists of Real Investment and Real Housing Investment. The third block consists of Bond Yield 10y and Commodity Prices. The fourth block consists of Loans NFCs and Oil Prices. The fifth block consists of EuroStoxx. The sixth block consists of Real Consumption. The seventh block consists of Effective Exchange Rate and US Real GDP. The eighth block consists of House Prices. The ninth block consists of M3 and Vstoxx.

³⁴Note also that we do not find support for a small VAR, where by “a small VAR” we mean the VAR with only the three variables of interest in the first block. The posterior odds in favor of ω^* relative to the best “small VAR” are approximately 3×10^6 .

³⁵We also try other ways of tightening the prior. We consider the model with the same block-exogeneity restrictions as the best model ω^* and we tighten the Sims-Zha prior by raising the hyperparameter on

How certain are we that a given variable is to be left out? Table 3 reports, for each remaining variable in the dataset, the best model in the family Ω with that variable in the first block. In particular, the posterior odds relative to the best model ω^* are given. As Kass and Raftery (1995), we think of the posterior odds between 1 and 0.3 as “not worth more than a bare mention”, of the odds between 0.3 and 0.05 as “positive”, of the odds between 0.05 and 0.007 as “strong”, and of the odds below 0.007 as “very strong”.

How strong is the evidence that a given variable is to be left out of the VAR model with the price level, GDP, and the policy rate? Consider Table 3. The evidence that the long-term interest rate is to be left out is borderline between “not worth more than a bare mention” and “positive”. We noted before that the long-term interest rate is the only variable appearing in Table 1 other than the unemployment rate, the bond spread, the purchasing managers index, and the federal funds rate. Next, the evidence that the price of oil, the stock market index, housing investment, the commodity price index, consumption, and the stock market implied volatility index are to be left out is “positive”. The evidence that U.S. GDP, loans to non-financial corporations, and the exchange rate are to be left out is “strong”. The evidence that the house price index and money supply are to be left out is “very strong”.

The last column of Table 3 shows the remaining variables in the first block of each model in that table. In all models except one the first block includes the unemployment rate, the bond spread, the purchasing managers index, and the federal funds rate.³⁶ We see this finding as consistent with the main finding that the best VAR with the price level, GDP, and the policy rate in addition includes those four variables.

Single block-exogeneity restriction versus multiple block-exogeneity restrictions. The data strongly support imposing multiple block-exogeneity restrictions rather than the “one-unit-root” dummy gradually from 1 (the best model) to 10. The posterior odds relative to the unrestricted model fall to about 12. Furthermore, the posterior odds of the model ω^* (with the “one-unit-root” dummy hyperparameter set to 1) relative to the model with the same block-exogeneity restrictions and the “one-unit-root” dummy hyperparameter set to 10 are enormous, about 4×10^{14} . The effect of raising the hyperparameter on the “no-cointegration” dummy to 10 is similar to the effect of reducing the “overall tightness” hyperparameter to 0.005.

³⁶The only exception is the best model with House Prices in the first block.

variable	(no)	odds to ω^*	variables from y_2 in the first block
Unemployment Rate	6	1.00	6,8,18,20
Bond Spread BBB 7to10y	8	1.00	6,8,18,20
PMI	18	1.00	6,8,18,20
US Fed Funds Rate	20	1.00	6,8,18,20
Bond Yield 10y	7	0.30	6,7,8,18,20
Oil Price	12	0.17	6,8,12,18,20
Euro Stoxx	16	0.15	6,8,16,18,20
Real Housing Investment	15	0.12	6,8,15,18,20
Real Investment	5	0.09	5,6,8,18,20
Commodity Prices	13	0.07	6,7,8,13,18,20
Real Consumption	4	0.06	4,6,7,8,18,20
VStoxx	17	0.06	6,8,17,18,20
US Real GDP	19	0.03	6,7,8,12,18,19,20
Loans NFC	10	0.03	6,7,8,10,18,20
Effective Exchange Rate	11	0.02	6,8,11,18,20
House Prices	14	0.005	8,12,14,18
M3	9	0.002	4,6,7,8,9,18,20

Table 3: The best model for each variable.

than one block-exogeneity restriction. The posterior odds in favor of the best model ω^* relative to the model that only has a single block-exogeneity restriction between the first block in the model ω^* and the other variables in the dataset is approximately 2×10^7 . Furthermore, we searched for the best model in the restricted family of models $\Omega' \subset \Omega$ such that the family Ω' includes VARs with at most one block-exogeneity restriction. The posterior odds in favor of the model ω^* relative to the best model in Ω' are approximately 10^4 . It turns out that in each of the best 387 models in the family Ω' the first block includes the unemployment rate, the bond spread, the purchasing managers index, and the federal

funds rate. We find this result reassuring.³⁷

Evidence from subsamples. We repeated the entire analysis in subsamples. First, we dropped the last four quarters. The results are very similar to the ones described above. Second, we dropped the last six quarters. The results differ somewhat from the ones described above. The main difference is that the price of oil and commodity prices do well in addition to unemployment rate, the bond spread, the purchasing managers index, and the federal funds rate. The ranking of the other variables changes, but house prices and money supply continue to be the worst variables. Third, we split our baseline sample into two subsamples with the middle of 2007 as the cut-off point. We refer to the subsample with the data from 1999Q1 to 2007Q2 as “the calm subsample”. We refer to the subsample with the data from 2007Q3 to 2010Q4 as “the crisis subsample”.³⁸ In both of these subsamples, medium-size VARs do better than large and small VARs. In both of these subsamples, money supply, the index of house prices, loans to non-financial corporations, and the exchange rate do poorly. In the calm subsample, the following variables do well: the price of oil, the purchasing managers index, consumption, and U.S. GDP. In the crisis subsample, the following variables do well: the bond spread and the federal funds rate. Note that the price of oil, consumption, and U.S. GDP do not make it into the best model in the full sample. Furthermore, the unemployment rate fails to do well in each subsample separately. We are not surprised that the findings differ somewhat between the calm subsample and the crisis subsample, because the two subsamples are quite different from each other. In the future, it will be useful to redo this paper’s analysis with non-linear models in this particular sample. In some non-linear models, such as VARs with Markov-switching and VARs with stochastic volatility, the principle behind the choice of variables will be the same as the

³⁷The MC³ chains that we ran within the family Ω' quickly gravitate towards models where the two blocks have approximately the same number of variables. Such models have roughly the maximum number of zero restrictions possible in a model in the family Ω' .

³⁸When analyzing the calm subsample, we use the training sample 1989Q1-1998Q4. Adding this training sample improves the marginal likelihood in the calm subsample compared with using the Sims-Zha prior only. When analyzing the crisis subsample, we use the training sample 1999Q1-2007Q2. Adding this training sample improves the marginal likelihood in the crisis subsample compared with using the Sims-Zha prior only. Adding this training sample improves the marginal likelihood in the crisis subsample also compared with using the Sims-Zha prior and the training sample 1989Q1-2007Q2.

principle laid out in this paper. However, in a non-linear model computation of marginal likelihood will be more complex than shown in Section 3.3.

We think of the results reported in this section as an illustration of the methodology proposed in this paper. We do not want to argue that the results reported in this section settle once and for all which variables are useful when the interest is in modeling in a VAR the price level, GDP, and the policy rate in the euro area. The sample available to us is too short for that and structural breaks may cause different variables to be useful in different periods.

6 Conclusions

We show how to evaluate conveniently the marginal likelihood implied by a VAR with one block-exogeneity restriction and multiple block-exogeneity restrictions. One can use these results in Bayesian tests of block exogeneity – Granger causality – in VARs. We employ these results to guide the choice of variables to include in a VAR with given variables of interest. The question of the choice of variables arises in most applications of VARs, whether in forecasting or impulse response analysis. Typically, the choice of variables occurs informally. We do not want to argue that the choice of variables must occur formally, using the methodology of this paper, in each Bayesian VAR from now on. We do want to suggest that: (i) the choice of variables can occur formally in a straightforward way, and (ii) even when the choice of variables occurs informally, it is useful to know what formal procedure this informal choice is meant to mimic.

A Sims-Zha prior

The prior used in this paper consists of two components: (i) an initial prior formulated before seeing any data, and (ii) a training sample prior. See Sections 3.2 and 5.2. This appendix gives the details concerning the initial prior. See Section 5.2 concerning the training sample prior.

The initial prior is the prior proposed by Sims and Zha (1998) and consists of the following four components.

The first component is the modified Minnesota prior. The modified Minnesota prior is

$$p(\text{vec } B | \Sigma, \omega^U) = \mathcal{N} \left(\text{vec} \begin{pmatrix} I_N \\ 0_{K-N \times N} \end{pmatrix}, \Sigma \otimes WW' \right), \quad (23)$$

where W is a diagonal matrix of size $K \times K$ such that the diagonal entry corresponding to variable n and lag p equals $\lambda_1/(\hat{\sigma}_n p^{\lambda_2})$. The terms λ_1 , λ_2 , and $\hat{\sigma}_n$ are hyperparameters. Let $\bar{P} = (1, \dots, P)$ and $\hat{\sigma} = (\hat{\sigma}_1, \dots, \hat{\sigma}_N)$. Then

$$W^{-1} = \text{diag} \left(\lambda_1^{-1} \bar{P}^{\lambda_2} \otimes \hat{\sigma}, \lambda_3^{-1} \right),$$

where λ_3 is a hyperparameter associated with the constant term γ . We implement the modified Minnesota prior with the dummy observations

$$Y_{Litterman} = W^{-1} \begin{pmatrix} I_N \\ 0_{K-N \times N} \end{pmatrix}, \quad X_{Litterman} = W^{-1}.$$

We set $\hat{\sigma}$ equal to standard deviations of residuals from univariate autoregressive models with P lags fit to the individual series in the sample.

The second component of the initial prior is the one-unit-root prior. The one-unit-root prior is implemented with the single dummy observation

$$Y_{one-unit-root} = \lambda_4 \bar{y}, \quad X_{one-unit-root} = \lambda_4 (\bar{y}, \dots, \bar{y}, 1),$$

where λ_4 and \bar{y} are hyperparameters. We set $\bar{y} = (1/P) \sum_{t=0}^{P-1} y_{-t}$, the average of initial values of y .

The third component is the no-cointegration prior. The no-cointegration prior is implemented with the N dummy observations

$$Y_{no-cointegration} = \lambda_5 \text{diag}(\bar{y}), \quad X_{no-cointegration} = \lambda_5 (\text{diag}(\bar{y}), \dots, \text{diag}(\bar{y}), 0),$$

where λ_5 is a hyperparameter.

The fourth component of the initial prior is an inverted Wishart prior about Σ with mean $\text{diag}(\hat{\sigma}^2)$. This prior is

$$\begin{aligned} p(\Sigma|\omega^U) &= \mathcal{IW}(ZZ', \nu_0) \propto |\Sigma|^{-(\nu_0+N+1)/2} \exp\left(-\frac{1}{2} \text{tr}(ZZ'\Sigma^{-1})\right) \\ &= |\Sigma|^{-(\nu_0+1)/2} |\Sigma|^{-N/2} \exp\left(-\frac{1}{2} \text{tr}(Z' - 0B)'(Z' - 0B)\Sigma^{-1}\right), \end{aligned}$$

where $Z_{N \times N}$ and ν_0 are hyperparameters. This prior is proportional to a likelihood of N observations with $Y_\Sigma = Z'$ and $X_\Sigma = 0_{N \times K}$, multiplied by the factor $|\Sigma|^{-(\nu_0+1)/2}$. We set $Z = \sqrt{\nu_0 - N - 1} \text{diag}(\hat{\sigma}_i)$, which implies that the mean of this prior is

$$E(\Sigma) = \frac{ZZ'}{\nu_0 - N - 1} = \text{diag}(\hat{\sigma}_n^2).$$

We set $\nu_0 = K + N = N(P + 1) + 1$. The reason for this choice for the value of ν_0 is as follows. The inverted Wishart density is proper when $\nu_0 > N - 1$. The Sims-Zha prior is proper when $\nu_0 > K + N - 1$, because K degrees of freedom are “used up” by the normal density of B . Therefore, as a rule of thumb we use the next integer after $K + N - 1$ setting $\nu_0 = K + N$.

Collecting all dummy observations introduced here yields

$$Y_{SZ} = \begin{pmatrix} Y_{Litterman} \\ Y_{one-unit-root} \\ Y_{no-cointegration} \\ Y_\Sigma \end{pmatrix}, \quad \text{and} \quad X_{SZ} = \begin{pmatrix} X_{Litterman} \\ X_{one-unit-root} \\ X_{no-cointegration} \\ X_\Sigma \end{pmatrix}.$$

The matrices Y_{SZ} and X_{SZ} appear in expression (22).

We use the following values of the hyperparameters: $\lambda_1 = 0.1$, $\lambda_2 = 1$, $\lambda_3 = 2$, $\lambda_4 = 1$, and $\lambda_5 = 1$. Furthermore, we set the number of lags to 1, that is, $P = 1$. We explored the effect of different hyperparameter values and different values of P on the marginal likelihood of the data implied by the unrestricted VAR estimated on the training sample 1989Q1-1998Q4. We used a grid of values for each of the hyperparameters and we used a grid of values for the number of lags. We found that the above values of the hyperparameters and one lag yield the highest value of the marginal likelihood. The values of $\lambda_2, \lambda_3, \lambda_4$ and

λ_5 that we found to be optimal are the same as the values used in Sims and Zha (1998). Our preferred value of λ_1 is one-half of the value used by Sims and Zha (1998), which implies that our prior is tighter. This is consistent with the suggestion of Giannone et al. (2010) to use tighter priors as the number of variables in the VAR increases.

B MC³

This appendix gives the details concerning the MC³ algorithm that we use. See also Section 3.4.

Given a model $\omega \in \Omega$, we define the *neighborhood* of this model denoted $nbr(\omega)$. The neighborhood of a model ω is the set of all models that differ from the model ω by the position of only one variable in the pattern of block-exogeneity restrictions. For example, one variable that belongs to the first block in the model ω may instead belong to the second block in a model $\omega' \in nbr(\omega)$. In general, the position of a variable can differ in one of four possible ways: (i) the variable may join the previous block, (ii) the variable may join the next block, (iii) the variable may become a block on its own prior to its current block, (iv) the variable may become a block on its own posterior to its current block. The MC³ chain moves as follows. Suppose the chain is at a model ω . We attach equal probability to each model in $nbr(\omega)$ and randomly draw a candidate model ω' from $nbr(\omega)$. We accept this draw with probability

$$\min \left\{ 1, \frac{\#nbr(\omega)p(Y|\omega')}{\#nbr(\omega')p(Y|\omega)} \right\}$$

where $\#nbr(\omega)$ denotes the number of models in $nbr(\omega)$.

When sampling from the family Ω' we used the convergence criterion proposed in George and McCulloch (1997). We ran two chains of one million draws. (That is, we drew a candidate model one million times, but not all of these draws were accepted.) The first chain started from the unrestricted model. The first chain made 435222 moves and it visited 13091 models. We saved these 13091 models and, when running the second chain, we checked how often the second chain visited these same models. The second chain started from the model where y_1 is block-exogenous to all elements of y_2 . This block-exogeneity assumption yields the smallest possible first block, in contrast to the unrestricted model

which has the largest possible first block. Thus, the second chain started as far as possible from the first chain, in the sense the number of moves required for the second chain to get to the starting model of the first chain was at a maximum. The second chain stayed within the set of the 13091 models from the first chain 99.2% of the time. This suggests excellent convergence. The inference concerning best models from both chains was exactly the same.

When sampling from the family Ω , we began by examining the properties of the Monte Carlo described in Section 3.3. We verified that when a model has only one block-exogeneity restriction – so that we know the marginal likelihood analytically – the Monte Carlo converges to the true value. Next, we checked 20 randomly selected models with multiple block-exogeneity restrictions. With $M = 1000$ draws we recovered the marginal likelihood on average up to a factor of 2, which is an insignificant difference when it comes to comparing marginal likelihoods. With $M = 10$ draws we recovered the marginal likelihood on average up to a factor of 7, which is a significant difference. We concluded that, given a time constraint, we faced a trade-off between precision of the Monte Carlo (a large value of M) and length of the chain. We proceeded in two steps. We first ran an MC³ chain with one million draws using $M = 10$. We then collected the best 100,000 models visited by this chain and we recomputed the marginal likelihood implied by each of those models using $M = 1000$.

Satisfying formal convergence criteria when sampling from the family Ω is much more difficult compared with the family Ω' , for two reasons. First, the family Ω contains many more models than the family Ω' . Second, very many models in the family Ω differ little from one another. For example, consider two models with several blocks. Moving one variable from the last block to the last-but-one block typically produces only a negligible change in the marginal likelihood. Therefore, we took a pragmatic approach. We ran five independent chains of length one million with different starting points and we asked if the lessons we could draw from these chains were the same. We found that the lessons were the same. Section 5.3 reports results obtained by stacking all five chains together.

Consider detailed examples. In four out of the five chains the best model in that chain had the same first block as the model ω^* . The exception was that the best model in one chain included the long-term interest rate in the first block, in addition to the unemployment rate,

the bond spread, the purchasing managers index, and the federal funds rate. Furthermore, on average seven – and at least five – models out of the best ten models in each chain had the same first block as the model ω^* . In all of the remaining cases the variables from the first block in the model ω^* were all present and only one additional variable was present (in 78% of the cases this additional variable was the long-term interest rate).

In the calm subsample and in the crisis subsample the results across chains were somewhat less consistent than in the full sample. We conjecture that in shorter samples the differences in marginal likelihoods are smaller, and thus probability is distributed more evenly across models.

Our experience suggests that the search for the best model in the family Ω' can be performed quickly and reliably with MC³ even when the number of the remaining variables in the dataset, N_2 , is twice larger than in our study. The search for the best model in the family Ω takes several days, and therefore is feasible today so long as the number of remaining variables in the dataset, N_2 , is not much larger than in our study and will become more and more reliable in larger datasets as computers improve.

References

- Bańbura, M., Giannone, D., and Reichlin, L. (2010). Large Bayesian vector auto regressions. *Journal of Applied Econometrics*, 25(1):71–92.
- Bauwens, L., Lubrano, M., and Richard, J.-F. (1999). *Bayesian Inference in Dynamic Econometric Models*. Oxford University Press, Oxford.
- Chib, S. (1995). Marginal likelihood from the Gibbs output. *Journal of the American Statistical Association*, 90(432):1313–1321.
- Cushman, D. O. and Zha, T. (1997). Identifying monetary policy in a small open economy under flexible exchange rates. *Journal of Monetary Economics*, 39(3):433–448.
- Dickey, J. M. (1971). The weighted likelihood ratio, linear hypotheses on normal location parameters. *The Annals of Mathematical Statistics*, 42(1):204–223.
- George, E. I. and McCulloch, R. E. (1997). Approaches for Bayesian variable selection. *Statistica Sinica*, 7:339–373.
- George, E. I., Sun, D., and Ni, S. (2008). Bayesian stochastic search for VAR model restrictions. *Journal of Econometrics*, 142(1):553–580.
- Geweke, J. (2005). *Contemporary Bayesian Econometrics and Statistics*. John Wiley and Sons, Hoboken, New Jersey.
- Geweke, J. and Amisano, G. (2011). Optimal prediction pools. *Journal of Econometrics*, 164(1):130–141.
- Giannone, D., Lenza, M., and Primiceri, G. (2010). Prior selection for vector autoregressions. Discussion paper, Université Libre de Bruxelles, European Central Bank, and Northwestern University.
- Granger, C. W. J. (1969). Investigating causal relations by econometric models and cross-spectral methods. *Econometrica*, 37(3):424–38.
- Hamilton, J. D. (1994). *Time Series Analysis*. Princeton University Press, Princeton, New Jersey.

- Kass, R. E. and Raftery, A. E. (1995). Bayes factors. *Journal of the American Statistical Association*, 90(430):773–795.
- Maćkowiak, B. (2007). External shocks, U.S. monetary policy and macroeconomic fluctuations in emerging markets. *Journal of Monetary Economics*, 54(8):2512–2520.
- Madigan, D. and York, J. (1995). Bayesian graphical models for discrete data. *International Statistical Review*, 63(2):pp. 215–232.
- OEIS (2011). On-line encyclopedia of integer sequences: A000670. Available at <http://oeis.org>.
- Sims, C. (2010). Lecture notes: causal orderings and exogeneity. Available at <http://sims.princeton.ftp.edu/yftp/times09/addednotes/gcp09.pdf>.
- Sims, C. A. (1972). Money, income, and causality. *American Economic Review*, 62(4):540–52.
- Sims, C. A. (1980). Macroeconomics and reality. *Econometrica*, 48(1):1–48.
- Sims, C. A. and Zha, T. (1998). Bayesian methods for dynamic multivariate models. *International Economic Review*, 39(4):949–68.
- Waggoner, D. F. and Zha, T. (2003). A Gibbs sampler for structural vector autoregressions. *Journal of Economic Dynamics and Control*, 28(2):349–366.
- Zha, T. (1999). Block recursion and structural vector autoregressions. *Journal of Econometrics*, 90(2):291–316.